

# Shaders

Teaching Modern Computer Graphics

Anton Gerdelan  
gerdela@scss.tcd.ie

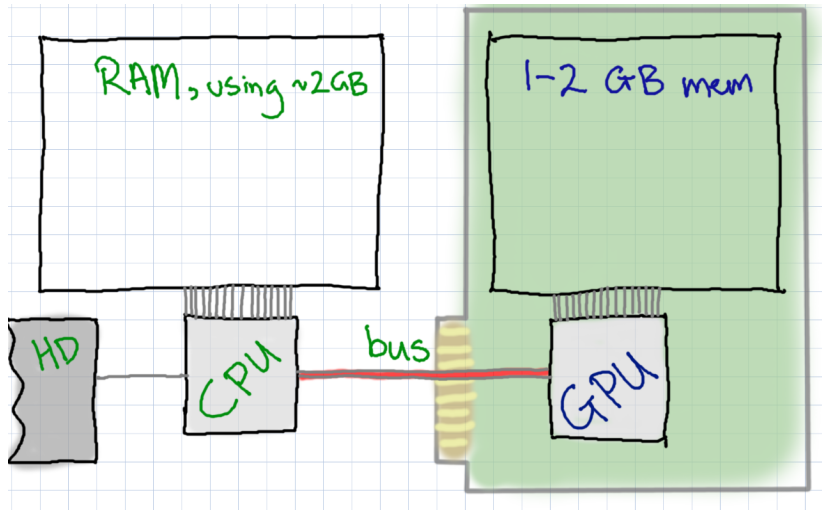
Knowledge and Data Engineering Group  
Trinity College Dublin, Ireland

March 5, 2014

# Overview

- ▶ Sharing some teaching and obsess...learning experiences
- ▶ Overview of the “new” computer graphics pipeline
- ▶ Shader programming
- ▶ WebGL and OpenGL ES
- ▶ Resources of note
- ▶ Some bold criticism / common problems

# Modern Hardware




# Data Parallelism

Copy all drawing data **before** drawing loop starts

- ▶ per-vertex “attribute” data points, normals, texture coordinates,...
- ▶ textures
- ▶ linked shader programmes
- ▶ buffers of commonly-used data (camera transformation matrices)

Update as seldom as possible e.g. camera matrices

# Graphics Pipeline



- ▶ Many new re-programmable stages
- ▶ Some optional
- ▶ Vertex Shader (transform) and Fragment Shader (colour) are basics

# APIs

In general:

- ▶ Direct3D - horrible OTT '90s-style object-oriented interface
- ▶ OpenGL - horrible crusty '80s-style state-machine C interface
- ▶ Must be good at memory allocation, pointers, addressing
- ▶ These are all huge problems for students learning
- ▶ Much better documentation for D3D
- ▶ Much better platform support for GL, docs are getting better-ish
- ▶ Functionality is almost 100% the same now (hardware-driven)
- ▶ GLSL and HLSL shaders are almost 100% the same
- ▶ Shaders are a lot of fun

# APIs

Latest versions:

- ▶ WebGL 1.0 / OpenGL ES 3.0 / OpenGL 2.1 / Direct3D 9 - vertex shaders + fragment shaders
- ▶ OpenGL 3.2 / Direct3D 10 - geometry shaders
- ▶ OpenGL 4 / Direct3D 11 - tessellation shaders
- ▶ OpenGL 4.3 / Direct3D 11 - compute shaders (GPGPU)

# What's Gone?

- ▶ `glBegin()` and `glEnd()` - gone
- ▶ `glMatrix` - gone
- ▶ `glLight` (point/spot/directional) - gone
- ▶ **glut** - plenty of modern alternatives
- ▶ You must know how the memory/processor architecture works (roughly)
- ▶ You must know how to do transformation matrices and dot products
- ▶ You must write the lighting and texturing algorithms by hand




## 'Client' side (what happens on CPU)

1. use supporting API to start an OS window/canvas
2. use API start a graphics 'context'
3. attach context to window/surface
4. load/download meshes, textures
5. use API to create buffer and copy meshes/textures onto graphics mem.
6. load shaders from strings. compile and link shaders. copy to GPU
7. switch shaders/geometry data to draw with
8. occasionally update a matrix
9. say `draw()` using current shaders and buffered data

In other words **not a lot** - just tie everything together

# “Hello Triangle”

Define a triangle of XYZ points in a vertex buffer



- ▶ can add an RGB colour for each point too
- ▶ just put into a arrays of 18 floats, total
- ▶ “bind” or enable this buffer before drawing

# “Hello Triangle”

## Vertex Shader in GLSL

```
#version 440

in vec3 vpos, vcolour;
out vec3 colour;

void main () {
    colour = vcolour;
    gl_Position = vec4 (vpos, 1.0);
}
```

- ▶ set each vertex position in homogeneous clip space
- ▶ output any variables to next stage in pipeline

# “Hello Triangle”

## Fragment Shader in GLSL


```
#version 440

in vec3 colour;
out vec4 frag_colour;

void main () {
    frag_colour = vec4 (colour, 1.0);
}
```

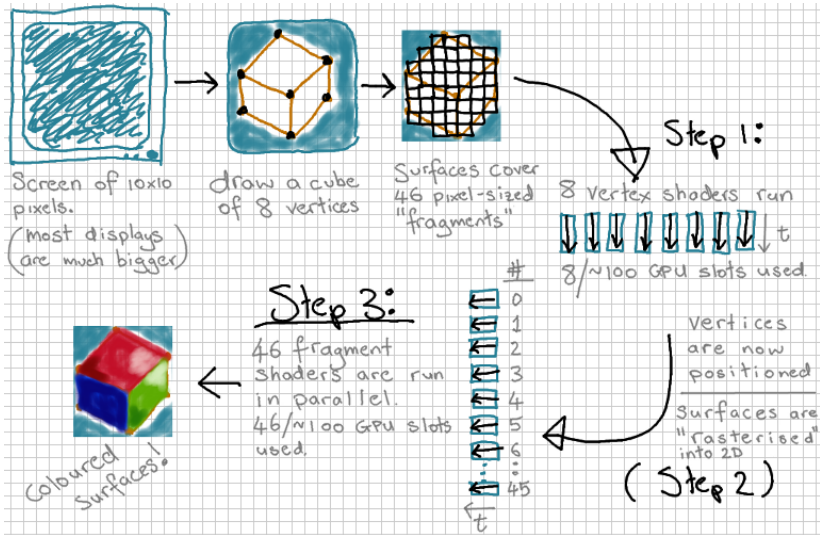
- ▶ set colour of each pixel-sized fragment on surface of geometry
- ▶ any input variables are **interpolated** to fragment position

# “Hello Triangle”




- ▶ Interpolation
- ▶ How many vertex shaders executed and how many fragment shaders executed?

# GPU Parallelism



# GPU Parallelism

- ▶ GPU driver knows how to optimise use of cores.
- ▶ Some drawing may start before other drawing has finished
- ▶ **not in sync with code running on CPU**
- ▶ Great tool called APITrace for visually profiling CPU and GPU calls <http://apitrace.github.io/>



- ▶ Can also do GPU performance queries in code
- ▶ Shaders are hard to debug. Mixed-arms approach

# GPU Hardware

## Selected High-End Desktop Adaptors

Year	Model	Shader Cores	Memory (GB)
2009	GeForce GTX 295	$2 \times 240$	1.6
2009	Radeon HD 5970	1600	2
2011	GeForce GTX 590	$2 \times 512$	3
2011	Radeon HD 6990	1536	4
2012	GeForce GTX 690	$2 \times 1536$	4
2013	Radeon HD 8990	2304	6
2014	GeForce GTX Titan Black	2880	6

Mobile adaptors follow a similar trend but 10 ~ 60% size ranges.



# Teaching Shaders

- ▶ most courses are still teaching **fixed-function** graphics
- ▶ '70s/'80s algorithms the same but practical side will never be used again
- ▶ **teach shaders**
- ▶ requires students build some knowledge of:
  - ▶ hardware (shader cores/bus/interpolation)
  - ▶ transformation pipeline and basic linear algebra
  - ▶ graphics algorithms to **write by-hand** (Phong, Catmull, etc.)
  - ▶ building and **linking libraries**
- ▶ requires relatively modern computers (but there are fall-backs)

# Choosing an API

- ▶ BTH - use any graphics API; Direct3D, OpenGL, MS XNA
- ▶ TCD - upgraded from pre-shaders GL to OpenGL 4
- ▶ Apple - 2.1, 3.2 core forward / 4.1 core forward
- ▶ Laptops - 2.1. Almost all have support for shaders / software emulation.
- ▶ fall back to 2.1, 3.2 - wider support, similar shaders, no tessellation
- ▶ **Android ADK/NDK is a technical nightmare world**
- ▶ iOS is okay but there is licence/fee/certificate nonsense
- ▶ or...

# WebGL

- ▶ OpenGL in a browser. Based on OpenGL ES 2.0
- ▶ Shaders are the same as ES, and GLSL 1.20 (OpenGL 2.1)
- ▶ **Very quick to develop**
- ▶ Runs on just about everything
- ▶ JavaScript 'glue' instead of C/C++
- ▶ No mess with libraries / IDEs
- ▶ Eric Haines' (Autodesk) course  
<https://www.udacity.com/course/cs291>
- ▶ Three.js library (super-easy option) [threejs.org](http://threejs.org)
- ▶ Experimenting with shaders [shadertoy.com](http://shadertoy.com)
- ▶ My 48-hour challenge Dolphin Rescue

# Problems

- ▶ Training staff - took me 2 years to get to a level of expertise with GL
- ▶ Lab support - programming basics, linking libraries, interface quirks
- ▶ *"It worked on my machine at home"*
- ▶ Supporting libraries for loading textures, opening windows, etc.
- ▶ Projects take a lot of student hours (think small or do groups)
- ▶ Will show embarrassing gaps in CS fundamentals (emergency C++ intro)
- ▶ Loading animated meshes

# Resources of Note



- ▶ Students do not like the go-to mega-tomes. Exorbitant, unclear.
- ▶ I put all my simplified GL4 teaching material on-line <http://antongerdelan.net/opengl>
- ▶ Maths cheat-sheet pdf and simple code